

# Analysis on Airbnb Dataset

Group2

11/13/2024

## R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
#Install Package
install.packages("readxl")
install.packages("tidyverse")
install.packages("leaflet")
install.packages("corrplot")

#Load libraries
library(readxl)
library(tidyverse) # This includes dplyr, ggplot2, and tidyr
library(tidyr)      # Explicitly load tidyr for pivot_longer
library(leaflet)    # Load leaflet for mapping
library(corrplot)   # Load corrplot for correlation visualization
```

## Load data into dataframe

```
# Read the Excel file
airbnb <- readxl::read_xlsx("data/AirbnbLA_2023.xlsx")

# View the data
head(airbnb) #use head for a sample
```

```
## # A tibble: 6 x 32
##       Id `Host Id` `Host Name`   `Host Is Superhost` `Host Acceptance Rate`
##   <dbl>   <dbl> <chr>           <lgl>                <chr>
## 1    109     521 Paolo             FALSE                50%
## 2   2708    3008 Chas.              TRUE                 100%
## 3   2732    3041 Yoga Priestess FALSE                42%
## 4  63416   309512 Vincenzo          TRUE                 96%
## 5  67089   210344 Brenna          TRUE                 95%
## 6   5728    9171 Sanni             FALSE                79%
## # i 27 more variables: `Host Response Rate` <chr>, `Host Response Time` <chr>,
## #   `Host Since` <dtm>, `Neighbourhood Group` <chr>, Neighbourhood <chr>,
## #   Latitude <dbl>, Longitude <dbl>, `Room Type` <chr>, Accommodates <dbl>,
## #   Beds <dbl>, Price <dbl>, `Instant Bookable` <lgl>, `First Review` <dtm>,
## #   `Last Review` <dtm>, License <chr>, `Reviews Per Month` <dbl>,
```

```
## # `Minimum Nights` <dbl>, `Number Of Reviews` <dbl>,
## # `Number Of Reviews L30D` <dbl>, `Number Of Reviews Ltm` <dbl>, ...
```

## Perform Data Cleaning

```
# Initial data cleaning and renaming columns
airbnb_cleaned <- airbnb %>%
  rename('id' = 'Id',
        'host_id' = 'Host Id',
        'host_name' = 'Host Name',
        'host_is_superhost' = 'Host Is Superhost',
        'host_acceptance_rate' = 'Host Acceptance Rate',
        'host_response_rate' = 'Host Response Rate',
        'host_response_time' = 'Host Response Time',
        'host_since' = 'Host Since',
        'neighbourhood_group' = 'Neighbourhood Group',
        'neighbourhood' = 'Neighbourhood',
        'latitude' = 'Latitude',
        'longitude' = 'Longitude',
        'room_type' = 'Room Type',
        'accommodates' = 'Accommodates',
        'beds' = 'Beds',
        'price' = 'Price',
        'instant_bookable' = 'Instant Bookable',
        'first_review' = 'First Review',
        'last_review' = 'Last Review',
        'license' = 'License',
        'reviews_per_month' = 'Reviews Per Month',
        'minimum_nights' = 'Minimum Nights',
        'number_of_reviews' = 'Number Of Reviews',
        'number_of_reviews_l30d' = 'Number Of Reviews L30D',
        'number_of_reviews_ltm' = 'Number Of Reviews Ltm',
        'review_scores_rating' = 'Review Scores Rating',
        'review_scores_accuracy' = 'Review Scores Accuracy',
        'review_scores_checkin' = 'Review Scores Checkin',
        'review_scores_cleanliness' = 'Review Scores Cleanliness',
        'review_scores_communication' = 'Review Scores Communication',
        'review_scores_location' = 'Review Scores Location',
        'review_scores_value' = 'Review Scores Value'
  )

# drop licence column
airbnb_cleaned <- select(airbnb_cleaned, -license)

# Convert "N/A" values to NA
airbnb_cleaned <- airbnb_cleaned %>%
  mutate(
    host_acceptance_rate = if_else(host_acceptance_rate == "N/A", NA, host_acceptance_rate),
    host_response_rate = if_else(host_response_rate == "N/A", NA, host_response_rate),
    host_response_time = if_else(host_response_time == "N/A", NA, host_response_time)
  )

# Impute missing review scores with the mean value of each column
airbnb_cleaned <- airbnb_cleaned %>%
```

```
mutate(
  review_scores_accuracy = if_else(is.na(review_scores_accuracy), mean(review_scores_accuracy, na.rm = TRUE), review_scores_accuracy),
  review_scores_checkin = if_else(is.na(review_scores_checkin), mean(review_scores_checkin, na.rm = TRUE), review_scores_checkin),
  review_scores_cleanliness = if_else(is.na(review_scores_cleanliness), mean(review_scores_cleanliness, na.rm = TRUE), review_scores_cleanliness),
  review_scores_communication = if_else(is.na(review_scores_communication), mean(review_scores_communication, na.rm = TRUE), review_scores_communication),
  review_scores_location = if_else(is.na(review_scores_location), mean(review_scores_location, na.rm = TRUE), review_scores_location),
  review_scores_value = if_else(is.na(review_scores_value), mean(review_scores_value, na.rm = TRUE), review_scores_value)
)

# Check for missing values again
colSums(is.na(airbnb_cleaned))
```

```
##           id           host_id
##           0             0
##    host_name    host_is_superhost
##           0             0
##    host_acceptance_rate    host_response_rate
##          4903             6076
##    host_response_time    host_since
##          6076             0
##    neighbourhood_group    neighbourhood
##           0             0
##          latitude    longitude
##           0             0
##          room_type    accommodates
##           0             0
##           beds    price
##           0             0
##    instant_bookable    first_review
##           0             0
##          last_review    reviews_per_month
##           0             0
##    minimum_nights    number_of_reviews
##           0             0
##    number_of_reviews_l30d    number_of_reviews_ltm
##           0             0
##    review_scores_rating    review_scores_accuracy
##           0             0
##    review_scores_checkin    review_scores_cleanliness
##           0             0
##    review_scores_communication    review_scores_location
##           0             0
##    review_scores_value
##           0
```

```
sum(is.na(airbnb_cleaned))
```

```
## [1] 17055
```

**Question 1: Which type of Airbnb properties garner the most reviews, indicating popularity?**

```
#Summarise the total review, avg rating, no. of reviews across room type
```


```

airbnb_popularity <- airbnb_cleaned %>%
  group_by(room_type) %>%
  summarise(total_reviews = sum(number_of_reviews),
            avg_rating = sum(review_scores_rating * number_of_reviews) / sum(number_of_reviews))

# Bubble chart for total reviews and avg_rating

ggplot(airbnb_popularity, aes(x = total_reviews, y = avg_rating, size = total_reviews, fill = room_type)) +
  geom_point(shape = 21, color = "black", stroke = 1) +
  geom_text(aes(label = paste(total_reviews, "\n", round(avg_rating, 2))),
            vjust = 2, color = "black", size = 5) +
  scale_size(range = c(3, 15), guide = "none") +
  scale_fill_brewer(palette = "Set1") +
  labs(title = "Popularity of Airbnb Room Types: Reviews vs. Ratings",
       x = "Total Reviews",
       y = "Average Rating",
       fill = "Room Type") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

```



**Question 2:** How does the distribution of listing types vary across different neighborhoods or regions?

```


airbnb_summary <- airbnb_cleaned %>%
  group_by(neighbourhood_group, room_type) %>%
  summarise(count = n(), .groups='drop') %>%

```

```
mutate(percentage = count / sum(count) * 100)

# Plot the 100% stacked bar chart for distribution of listing types vary across different neighborhoods

ggplot(airbnb_summary, aes(x = neighbourhood_group, y = percentage, fill = room_type)) +
  geom_bar(stat = "identity", position = "fill") +
  scale_y_continuous(labels = scales::percent) +
  labs(x = "Neighbourhood", y = "Percentage", fill = "Room Type",
       title = "Distribution of Listing Types by Neighbourhood") +
  theme_minimal()
```



**Question 3: Who are the top 10 Super hosts based on listings, review scores, and number of reviews? How do their listing and review score distributions vary?**

```
#Assign the rank to each host based on review score rating, number of reviews and total listings
airbnb_groupby_unique_host <- airbnb_cleaned %>%
  filter(host_is_superhost = TRUE) %>%
  group_by(host_id, host_name) %>%
  summarise(avg_review_score = mean(review_scores_rating),
            total_reviews = sum(number_of_reviews),
            total_listing = n()) %>%
  ungroup() %>% # Corrected to 'ungroup()'
  mutate(
    rank_review_score = rank(-avg_review_score), # Rank by average review score rating
    rank_number_of_reviews = rank(-total_reviews), # Rank by number of reviews
    rank_number_of_listings = rank(-total_listing) # Rank by maximum number of listings
```

```


)

# Combine the ranks (e.g., by summing them up for an overall rank)
filtered_data <- airbnb_groupby_unique_host %>%
  mutate(overall_rank = rank_review_score + rank_number_of_reviews + rank_number_of_listings) %>%
  arrange(overall_rank) # Sort by the combined rank

# Select the top 10 super hosts
top_10_hosts <- filtered_data %>%
  slice_head(n = 10) %>%
  select(host_id, host_name, avg_review_score, total_reviews, total_listing, overall_rank)

# Visualize top 10 Super hosts based on listings, review scores, and number of reviews
ggplot(top_10_hosts, aes(x = total_reviews, y = avg_review_score, color = host_name, size = total_listing)) +
  geom_point(alpha = 0.7) +
  scale_size_continuous(range = c(3, 10)) +
  labs(title = "Top 10 Super Hosts by Average Review Score and Total Reviews",
       x = "Total Number of Reviews",
       y = "Average Review Score",
       size = "Total Listings",
       color = "Super Host Name") +
  theme_minimal()

```



## Question 4: What is the overall price trend for different room types on Airbnb?

```
# Calculate summary statistics for room types and see for price outliers

summary_stats_by_roomType <- airbnb_cleaned %>%
  group_by(room_type) %>%
  summarise(
    Average_Price = mean(price, na.rm = TRUE),
    Median_Price = median(price, na.rm = TRUE),
    Min_Price = min(price, na.rm = TRUE),
    Max_Price = max(price, na.rm = TRUE),
    SD_Price = sd(price, na.rm = TRUE)
  )

print(summary_stats_by_roomType)


## # A tibble: 4 x 6
##   room_type      Average_Price Median_Price Min_Price Max_Price SD_Price
##   <chr>          <dbl>         <dbl>    <dbl>    <dbl>    <dbl>
## 1 Entire home/apt      268.           170         5     99999     777.
## 2 Hotel room           798.           100        22      9999    2439.
## 3 Private room         118.            69        10     99999    1204.
## 4 Shared room          53.7            35        12     1200     95.3

# Data Cleaning: Remove outliers in price using the IQR method
remove_price_outliers <- function(data) {
  Q1 <- quantile(data$price, 0.25, na.rm = TRUE)
  Q3 <- quantile(data$price, 0.75, na.rm = TRUE)
  IQR_value <- Q3 - Q1
  lower_bound <- Q1 - 1.5 * IQR_value
  upper_bound <- Q3 + 1.5 * IQR_value
  data %>% filter(price >= lower_bound & price <= upper_bound)
}

airbnb_filtered <- remove_price_outliers(airbnb_cleaned)
```

## Plot Analysis Question4

```
# Box plot of Price Distribution by Room Type
ggplot(airbnb_filtered, aes(x = room_type, y = price, fill = room_type)) +
  geom_boxplot() +
  stat_summary(fun = "mean", geom = "point", shape = 18, size = 3, color = "red", fill = "red",
    position = position_dodge(width = 0.75)) + # Adjusts the position of the mean marker
  labs(title = "Price Distribution by Room Type",
    x = "Room Type",
    y = "Price ($)",
    fill = "Room Type") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



**Question 5: How does the average price of Airbnb listings vary across different neighborhoods in Los Angeles?**

```
# Group the data by neighborhood and calculate average price
df_grouped <- airbnb_filtered %>%
  group_by(neighbourhood) %>%
  summarise(
    Avg_Price = mean(price, na.rm = TRUE),
    latitude = first(latitude),
    longitude = first(longitude),
    .groups = 'drop'
  )

print(df_grouped)
```

```
## # A tibble: 265 x 4
##   neighbourhood Avg_Price latitude longitude
##   <chr>          <dbl>    <dbl>    <dbl>
## 1 Acton          171.      34.5     -118.
## 2 Adams-Normandie  90.5      34.0     -118.
## 3 Agoura Hills    174.      34.2     -119.
## 4 Agua Dulce      158.      34.5     -118.
## 5 Alhambra        128.      34.1     -118.
## 6 Alondra Park    178.      33.9     -118.
## 7 Altadena        155.      34.2     -118.
## 8 Angeles Crest   168.      34.4     -118.
```



```
## 9 Arcadia          123.      34.1    -118.
## 10 Arleta          100       34.2    -118.
## # i 255 more rows

# Create a color palette based on average prices
pal <- colorNumeric(palette = "viridis", domain = df_grouped$Avg_Price)

# Create the interactive map with color tones
leaflet(df_grouped) %>%
  addTiles() %>%
  addCircleMarkers(
    lng = ~longitude,
    lat = ~latitude,
    radius = ~Avg_Price / 50,
    popup = ~paste(neighbourhood, ": $", round(Avg_Price, 2)),
    color = ~pal(Avg_Price),
    fillOpacity = 0.7
  ) %>%
  setView(lng = mean(df_grouped$longitude), lat = mean(df_grouped$latitude), zoom = 11) %>%
  addLegend("bottomright", pal = pal, values = ~Avg_Price,
    title = "Average Price",
    opacity = 0.7)
```

## Data Modeling Visualaization

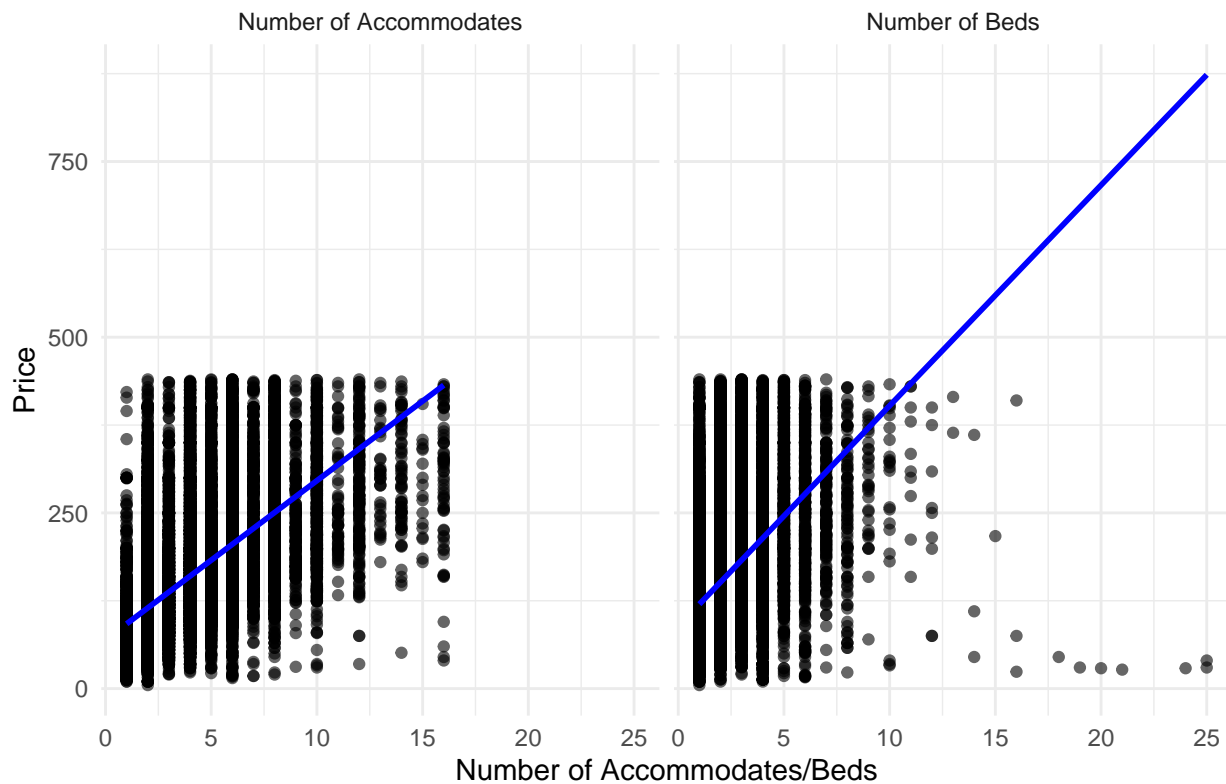
```
# Impact of Beds and Accommodates on Price of Room Types
# Checking which has more impact: Accommodates or Beds
correlation_matrix <- airbnb_filtered %>%
  select(price, accommodates, beds) %>%
  cor()
print(correlation_matrix)

##              price accommodates      beds
## price      1.0000000    0.6022625 0.4964012
## accommodates 0.6022625    1.0000000 0.8219663
## beds       0.4964012    0.8219663 1.0000000

# Reshape the data for combined plotting
airbnb_long <- airbnb_filtered %>%
  pivot_longer(cols = c(beds, accommodates), names_to = "Type", values_to = "Value")

# Create a single graph for Price vs. Beds and Price vs. Accommodates
ggplot(airbnb_long, aes(x = Value, y = price)) +
  geom_point(alpha = 0.6) + # Adjust transparency for better visibility
  geom_smooth(method = "lm", se = FALSE, color = "blue") + # Add linear regression line
  labs(title = "Price vs. Accommodates and Beds",
    x = "Number of Accommodates/Beds",
    y = "Price") +
  facet_wrap(~ Type, labeller = as_labeller(c(beds = "Number of Beds", accommodates = "Number of Accommodates")))
theme_minimal()
```

## Price vs. Accommodates and Beds




Question 6: Is there a correlation between the number of reviews and overall ratings? Do hosts with more reviews tend to have better ratings?

```
# Transform data for faceting
airbnb_long <- airbnb_cleaned %>%
  pivot_longer(cols = c(number_of_reviews, number_of_reviews_l30d, number_of_reviews_ltm),
               names_to = "review_type",
               values_to = "review_count")

# Creating a scatter plot with a trend line
ggplot(airbnb_long, aes(x = review_scores_rating, y = review_count)) +
  geom_smooth(aes(color = review_type), se=FALSE) +
  labs(title = "Correlation between Number of Reviews and Overall Rating",
       x = "Overall Rating (Review Scores)",
       y = "Number of Reviews") +
  theme_minimal()
```

## Correlation between Number of Reviews and Overall Rating



Question 7: Which factors, such as the check-in process, cleanliness, accuracy of listing descriptions, etc., most significantly impact review ratings?

```
cor_matrix <- cor(airbnb_cleaned[, c("review_scores_rating", "review_scores_accuracy",
                                     "review_scores_cleanliness", "review_scores_communication",
                                     "review_scores_checkin", "review_scores_value",
                                     "review_scores_location")],
                 use = "complete.obs")
```

```
print("Correlation matrix of factors impacting review ratings:")
```

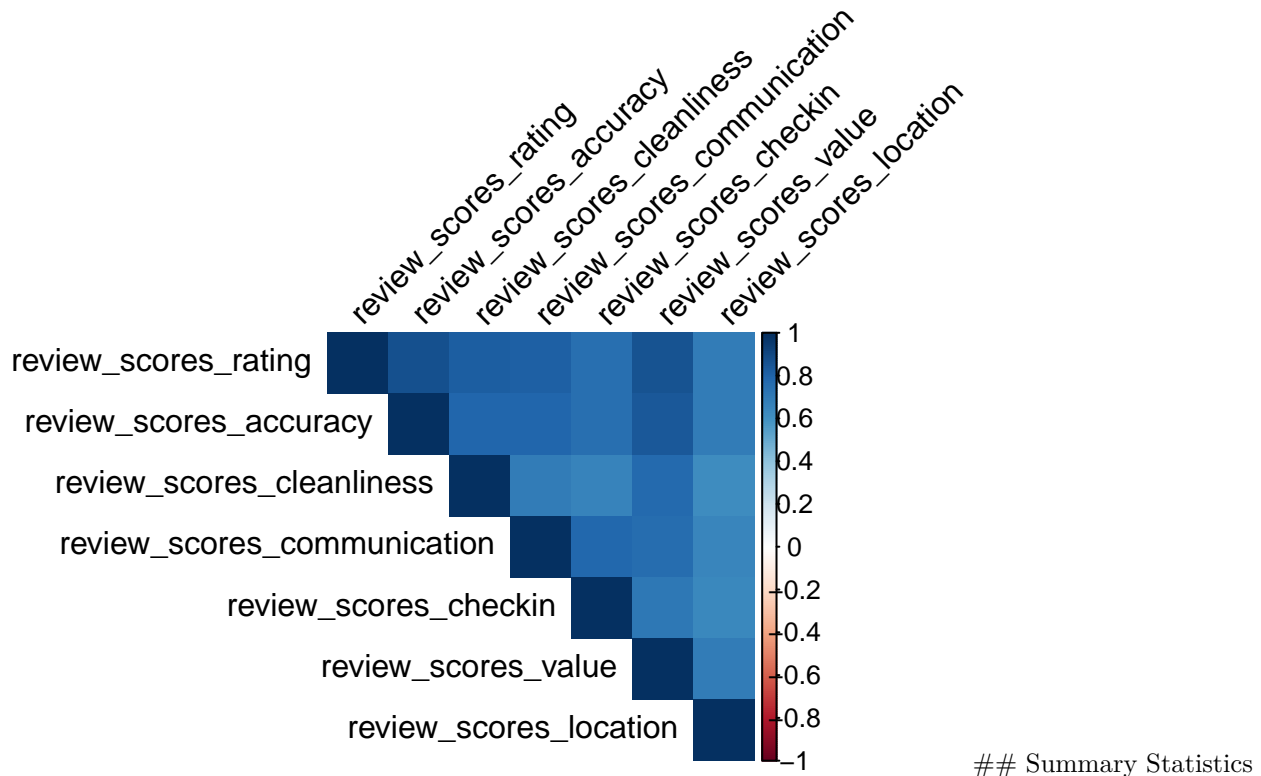
```
## [1] "Correlation matrix of factors impacting review ratings:"
```

```
print(cor_matrix)
```

```
##               review_scores_rating review_scores_accuracy
## review_scores_rating              1.0000000             0.8731444
## review_scores_accuracy              0.8731444             1.0000000
## review_scores_cleanliness           0.8281561             0.7925776
## review_scores_communication         0.8103266             0.7936822
## review_scores_checkin               0.7569332             0.7532662
## review_scores_value                 0.8691172             0.8431411
## review_scores_location              0.6952753             0.6936656
##               review_scores_cleanliness
## review_scores_rating              0.8281561
## review_scores_accuracy            0.7925776
## review_scores_cleanliness         1.0000000
```

```
## review_scores_communication 0.6999623
## review_scores_checkin 0.6647535
## review_scores_value 0.7751937
## review_scores_location 0.6201472
##
## review_scores_communication review_scores_checkin
## review_scores_rating 0.8103266 0.7569332
## review_scores_accuracy 0.7936822 0.7532662
## review_scores_cleanliness 0.6999623 0.6647535
## review_scores_communication 1.0000000 0.7836596
## review_scores_checkin 0.7836596 1.0000000
## review_scores_value 0.7681240 0.7180183
## review_scores_location 0.6556051 0.6442605
##
## review_scores_value review_scores_location
## review_scores_rating 0.8691172 0.6952753
## review_scores_accuracy 0.8431411 0.6936656
## review_scores_cleanliness 0.7751937 0.6201472
## review_scores_communication 0.7681240 0.6556051
## review_scores_checkin 0.7180183 0.6442605
## review_scores_value 1.0000000 0.6993991
## review_scores_location 0.6993991 1.0000000

# Visualize correlations
corrplot(cor_matrix, method = "shade", type = "upper", tl.col = "black", tl.srt = 45)
```



```
#Summary Statistics for Listing Capacity across Room Type

get_mode <- function(x) {
  uniqx <- unique(x)
  uniqx[which.max(tabulate(match(x, uniqx)))]
}
```

```
summary_stats_accommodates <- airbnb_cleaned %>%
  group_by(room_type) %>%
  summarize(
    mean = mean(accommodates, na.rm = TRUE),
    median = median(accommodates, na.rm = TRUE),
    min = min(accommodates, na.rm = TRUE),
    max = max(accommodates, na.rm = TRUE),
    sd = sd(accommodates, na.rm = TRUE),
    mode = get_mode(accommodates),
    Inter_quertile = IQR(accommodates, na.rm=TRUE),
    count = n())
```

```
summary_stats_accommodates
```

```
## # A tibble: 4 x 9
##   room_type      mean median   min   max    sd  mode Inter_quertile count
##   <chr>      <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl>      <dbl> <int>
## 1 Entire home/apt 4.66      4     1    16 2.81     2         4 24493
## 2 Hotel room     2.34      2     1     6 0.922    2         0   62
## 3 Private room   1.99      2     1    16 1.11     2         1  7530
## 4 Shared room    2.24      1     1    16 2.33     1         1   364
```

Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing of the R code that generated the plot.